

## ECS 158 Project 3

Sean Malloy, 998853013

Joseph,

ECS 158, Joël Porquet

11/18/2018

### Program 1: Matrix-multiplication with message passing

We decided to implement matrix multiplication in `multiply_mpi` by using the `ikj` matrix multiplication from Project 1. The main function handles input validation and initializes the MPI environment. We separated the work by a master-slave working paradigm with `com0` being the master and the rest of the created processor `coms` being the slaves. Initially the master will allocate and initialize the matrices A, B and C. The master will then send the relevant chunk sizes, parts of the matrix A and also send B via `MPI_Send` and `MPI_Bcast`. We sent the whole B matrix since our Matrix Multiplications required it to work correctly, and we utilized `MPI_Bcast` to deliver the data since this seemed to be faster than using the `MPI_Send` method in certain situations, saving as much as 40 seconds in some runtimes. After the master sends the data, the slaves will receive these parameters via `MPI_Recv/MPI_Bcast`, and then continue to start working on their own pieces. The slaves will calculate the relevant parts of the C matrix independently. They will all use the `multiply_mpi` function to calculate a piece of C and send the calculated piece back to the master via `MPI_Send`. The master will be listening for these messages using `MPI_Recv` and an offset variable will keep track of which parts of the C matrix are being received to rebuild the C matrix.

Locally						
N	8 Processors	4 Processors	2 Processors	Serial	OMP	Pthread
500	0.046298	0.039501	0.034915	0.076914	0.042465	0.020875
	0.045176	0.039078	0.034759	0.080405	0.019945	0.022151
	0.045816	0.04053	0.034579	0.072048	0.029095	0.021396
Average	0.045763333	0.039703	0.034751	0.076455667	0.030501667	0.021474
1000	0.416857	0.394218	0.422085	0.577559	0.14836	0.14603
	0.39572	0.42886	0.420755	0.578746	0.220676	0.143488
	0.450489	0.430156	0.422429	0.587785	0.238844	0.141652
Average	0.421022	0.41774467	0.421756333	0.581363333	0.202626667	0.143723333
2000	3.136299	3.476246	3.338978	5.477773	1.310048	1.159077
	3.252358	3.309672	3.338753	5.469149	1.564502	1.158264
	3.098507	3.48029	3.338762	5.485127	1.591594	1.173145
Average	3.162388	3.42206933	3.338831	5.477349667	1.488714667	1.163495333
3000	10.477089	11.670172	11.200275	19.187115	3.86367	3.940683
	10.70196	11.392818	11.164662	19.175987	3.997528	3.891222
	10.463212	11.585681	11.204879	19.192328	4.011433	4.026503
Average	10.54742033	11.549557	11.18993867	19.18514333	3.957543667	3.952802667

For the local times for this implementation, we can see that if we increased the number of processors for MPI then it would be slower for smaller matrix sizes due to overhead compared to runs with less processors but will become much faster in larger matrix sizes since it can parallelize the data better to get it done much faster. Compared to OMP and Pthread MPI seems to be a bit slower, but that's because of the way that MPI must parallelize its data creating more overhead for its functions.

Cluster					
N	16 Processors	8 Processors	4 Processors	2 Processors	Serial
500	0.460313	0.766185	0.033531	0.037147	0.062168
	0.462432	0.790398	0.027087	0.036165	0.059812
	0.46448	0.78993	0.036316	0.036225	0.062659
Average	0.462408333	0.782171	0.032311333	0.036512333	0.061546333
1000	1.948763	1.565718	0.346518	0.348843	0.588272
	1.946731	1.573371	0.361926	0.376295	0.590393
	1.939037	1.578041	0.367589	0.394079	0.587093
Average	1.944843667	1.57237667	0.358677667	0.373072333	0.588586
2000	8.973455	6.99989	2.840205	3.065135	4.834275
	8.948229	6.988381	2.867304	3.06897	4.827524
	8.97409	6.995889	2.852408	3.06989	4.818571
Average	8.965258	6.99472	2.853305667	3.067998333	4.82679
3000	21.055664	17.302096	9.535181	10.436245	16.811883
	21.044934	17.272325	9.507568	10.448655	16.930845
	20.999419	17.273446	9.485279	10.423371	16.871649
Average	21.033339	17.2826223	9.509342667	10.43609033	16.871459
5000	63.151831	56.784605	43.748199	47.881802	77.79539
	62.639295	56.687771	43.608063	48.062257	78.055581
	62.718172	56.86574	43.699882	47.966169	77.667931
Average	62.83643267	56.779372	43.68538133	47.970076	77.839634

For the timing of the MPI\_Matrix multiplication, when we parallelize the data over a cluster, then it seems that the larger number of processors carry more overhead, as it needs to wait to communicate with the other computers over the network, and thus fewer processors might be better in some situations. We can see that having it done serially is usually slower than doing it parallelized, unless if we have to communicate with multiple computers to release the data to.

## Program 2: Mandelbrot

We use the same master-slave working paradigm as in Program 1 for Program 2, with com0 being the master and coms being the slaves. The main function again handles the input validation and initialization of the MPI environment. The master starts out by calculating the

relevant chunk sizes and starting indexes and sends the parameters that the slaves will be using with a combination of MPI\_Send and MPI\_Broadcast (Bcast). (chunk\_size, start, xcenter, ycenter, cutoff, increment). The slaves will then receive all the parameters via a combination of MPI\_Recv and MPI\_Broadcast and start working on the chunk they have received.

The function compute\_mandelbrot will map the complex plane with the 1024x1024 matrix. We iterate over the 1024x1024 matrix and align the index of the image of the mandelbrot set with the matrix. In order to map the correct image, we had to start iterating from the bottom left corner of the image. We calculate the x distance from the center to the left side of the image using xcenter and the increment and since we know that it is being mapped to a 1024x1024 matrix the distance would equate to 512\*increment. We do the same for the y distance and with the new x and y coordinates we have found the bottom left corner. Using the start and chunk sizes we are able to increment our x's and y's independently in order to calculate whether values are in mandelbrot set or not.  $z_{next} = z_{curr} + z_{curr} * C$  with  $C = xvalue + yvalue * I$  and  $|z_{next}| > 2$ . Once the mandelbrot set is calculated for the chunk the chunk is returned back to the master. The master rebuilds the 1024x1024 matrix from the slaves using MPI\_Recv. The matrix is transposed and written to a file to give the output mandelbrot pgm file.

MandelBrot						
Cutoff	32 Processors	16 Processors	8 Processors	4 Processors	2 Processors	Serial
100	0.548588	0.372809	0.220856	0.150223	0.173559	0.18545
	0.7558	0.382379	0.222744	0.145613	0.176846	0.184389
	0.700993	0.369012	0.218491	0.167863	0.175918	0.181043
Average	0.668460333	0.37473333	0.220697	0.154566333	0.175441	0.183627333
300	0.686276	0.426595	0.361097	0.407893	0.450675	0.461274
	0.749883	0.400628	0.32376	0.409385	0.457917	0.458505
	0.836359	0.369147	0.34272	0.386961	0.450339	0.452632
Average	0.757506	0.39879	0.342525667	0.401413	0.452977	0.457470333
500	1.088704	0.547712	0.490293	0.658137	0.7263	0.742486
	0.982334	0.485668	0.525751	0.654951	0.725829	0.728372
	0.937291	0.511387	0.499431	0.672994	0.724399	0.7282
Average	1.002776333	0.51492233	0.505158333	0.662027333	0.725509333	0.733019333
800	1.282963	0.758324	0.769739	1.025645	1.134734	1.133429
	1.473108	0.756142	0.736083	1.032585	1.128297	1.135651
	1.409152	0.747725	0.739019	1.029471	1.147957	1.135848
Average	1.388407667	0.75406367	0.748280333	1.029233667	1.136996	1.134976

As for the times, we had utilized a xcenter = -1, y center= -1, and zoom = 0.1 and differing cutoff values since that seemed the most optimal way to check the difference in speedup. From what we saw, if we had too many processors, then the speed will slow down due to overhead, but with about 8 processors we see the greatest speedup except for a small cutoff point. We also

see Serial being pretty decent for a small cutoff point, but get slower compared to our 8 processor implementation with larger cutoffs.